

ORADIA: Optimizing Resource Allocations for Dynamic Interactive Applications

Authors Redacted

Abstract

rewrite duplicate sentences

Modern computing systems are under intense pressure to provide guaranteed responsiveness to their workloads. Ideally, applications with strict performance requirements should be given just enough resources to meet these requirements consistently, without unnecessarily siphoning resources from other applications. However, executing multiple parallel, real-time applications while satisfying QoS requirements is a complex optimization problem and traditionally operating systems have provided little support to provide QoS to applications. As a result, client, cloud and embedded systems have all resorted to over-provisioning and isolating applications to guarantee responsiveness. We present ORADIA, a resource allocation framework, which is designed to provide responsiveness guarantees to a simultaneous mix of high-throughput parallel, interactive, and real-time applications in an efficient, scalable manner. By measuring application behavior directly and using convex optimization techniques, ORADIA is able to understand the resource requirements of applications and perform near optimal resource allocation with low overhead.

1. Introduction

The demands on modern systems have changed in the Post-PC era. Users demand application responsiveness. These applications often deliver high-quality multimedia that require real-time guarantees. Meeting this responsiveness goal is a challenge for all types of systems including cloud systems, databases, web servers, client operating systems, and emerging distributed embedded systems. Additionally, battery life and system power are extremely important, thereby forcing systems to try to find more efficient ways to meet the needs of their workloads.

Applications, jobs, or queries with strict performance requirements should ideally be given just enough system resources (*e.g.* nodes, processor cores, cache slices, memory pages, various kinds of bandwidth) to meet these requirements consistently, without unnecessarily siphoning resources from other applications. However, executing multiple parallel, real-time applications while satisfying *Quality-of-Service* (QoS) requirements is a complex optimization problem, particularly as modern hardware diversifies to include a variety of parallel architectures (*e.g.*, multicore, gpus). Operating systems have historically not provided useful mechanisms that implement stronger performance guarantees and resource allocation has been rather unsystematic, making it difficult to reason about the expected response time of an application.

Consequently, responsiveness has been obtained at significant expense by designing for the worst case and over-

provisioning. Evidence of this behavior can be found in current systems of all sizes. OSs attempt to describe responsiveness with a single value (usually called a *priority*) associated with a thread of computation and adjusted within the operating system by a variety of *ad-hoc* mechanisms. Other shared resources either employ independent machinery (*e.g.* memory, caches), or are deemed so abundant as to require no explicit management at all (*e.g.* I/O, network bandwidth). Priority approaches have no mechanism to describe deadlines or the resources required to meet them and so must run the highest priority applications as fast as possible on all the resources requested. As a result, interactive and real-time applications are often run needlessly fast with significant over-provisioning — wasting power and energy and preventing other applications from using the resources.

Some mobile systems have gone so far as to limit which applications can run in the background [2] in order to preserve responsiveness and battery life, despite the obvious concerns this raises for user experience. Cloud computing providers routinely utilize their clusters at only 10% to 50% to keep the system responsive despite the additional operational costs of consuming electricity and the significant impact to the capital costs of the infrastructure [3, 28]. In some cases, clusters only run a single application on each cluster to avoid unexpected interference. Similarly the realtime community has used completely separate systems for each application to provide QoS despite the high cost of specialization and low utilization with this approach.

In this paper, we present ORADIA, a resource allocation framework designed to provide responsiveness guarantees to a simultaneous mix of high-throughput parallel, interactive, and real-time applications in an efficient, scalable manner. Unlike traditional systems, ORADIA considers all resource types when making decisions and continuously determines the complete set of resources each application will use. It maintains application-specific performance models through measurement to help determine resource requirements and responds to dynamic changes in both the deadlines and the relative importance of meeting each of them. It leverages convex optimization to determine the quantity of each resource to give each application. This lets the system make trade-offs among application QoS/responsiveness, system performance, and energy efficiency.

We believe ORADIA is applicable to many resource allocation scenarios including cloud providers determining how much to give each job to avoid violating SLAs, databases allocating resources to queries, and distributed embedded systems allocating bandwidth among devices and sensors. In this paper we choose to study client systems and their applications, im-

plementing ORADIA in a general-purpose operating system, because we believe this scenario has some of the most significant resource allocation challenges: a constantly changing environment requiring low overhead and fast response times, shared resources that create more interference among the applications, and platforms that are too diverse to allow *a priori* performance prediction.

Add Performance Numbers

2. ORADIA Architecture

ORADIA is a framework designed to determine the proper amount of each resource type (e.g., nodes, processor cores, cache slices, memory pages, various kinds of bandwidth) to assign to each application. For example, consider a scenario of a video conference, where each participant requires a separate, performance guaranteed video stream. New participants may join the conference and others leave, increasing or decreasing the number of streams running at any given time. Simultaneously, participants may be collaborating through browsers, watching shared video clips and web searching, while running compute-intensive background tasks such as updates, virus scans, or file indexing. Although it may be relatively straightforward to provide responsiveness guarantees for applications such as video streams, the real challenge is to do so without using static resource reservations that compromise system utilization. The purpose of ORADIA is to dynamically assign resources to applications to guarantee responsiveness without over-provisioning and to adapt these allocations as the applications in the system change.

For our purposes an application is an entity to which the system allocates resources: it can be viewed as a complete application (e.g., a video player), a component of an application (e.g., a music synthesizer) or a process (e.g., indexing) for operating systems, jobs for warehouse-scale computing, and distributed applications in distributed embedded systems.

ORADIA is designed to work in systems with hierarchical scheduling. In a client system, this design just looks like classic two-level scheduling. ORADIA makes allocation decisions about resources (e.g., cores and memory pages), and micro-management of the resources within an application is left to user-level runtimes such as the Intel Threaded Building Blocks runtime [16] or Lithe [52] and to memory managers. In a cloud environment, ORADIA allocates resources (e.g., nodes and storage) to jobs, and scheduling is left to other entities such as the MapReduce framework [18] or the operating system on the node. Separating resource allocation from scheduling enables the use of application-specific scheduling policies, which have the potential to be easier to design and more efficient than a general-purpose scheduler that has to work for everything. This approach leaves the system to focus on the problem of *how many* of each resource type to assign to each application.

ORADIA takes a different approach to resource allocation than traditional systems relying heavily on application-specific

functions built through measurement and convex optimization. ORADIA formulates resource allocation as an optimization problem built from two types of application-specific functions: a Response Time Function and a Penalty function. The Response Time Function represents the measured performance of the application on different resources. The Penalty Function represents the user-level goals for the application (i.e., the deadline and how important it is to make that deadline). The following sections describe each of the functions and the optimization construction.

Response Time Functions

Response Time Functions (RTF) represent the expected *response time* of an application as a function of the resources allocated to the application. The response time is an application-specific measure of the performance of the application. For example, the response time of an application might be:

- The time from a mouse click to its result;
- The time to produce a frame;
- The time from a service request to its response;
- The time from job launch to job completion;
- The time to execute a specified amount of work.

The RTFs are built to be a convex function. All applications have a function of the same form but the application-specific weights are set using the performance history of the application. RTFs are designed to capture information such as how well an application scales with a particular resource. As a result, RTFs naturally support heterogeneity. Each CPU or GPU type is simply viewed as a different resource type by the system, and thus the RTFs will represent how effectively an application uses a particular type of core. Figure 1 shows two example RTFs we have created from applications we studied.




Figure 1: Response Time Functions for a breadth-first search and a stencil algorithm with 2 response dimensions: cores and cache slices.

Equation 1 below shows the form of the RTF we use in ORADIA.

$$\tau(w, a) = \tau_0 + \sum_{i \in n, j \in n} \frac{w_{i,j}}{\sqrt{a_i * a_j}} \quad (1)$$

Here τ represents the response time, i and j represent resource types, n represents the total number of resource types, a_i and a_j represent the allocation of resources types i and j ,

and $w_{i,j}$ represents the application-specific weight for the term representing resource i * resource j . The dimensionality of the function increases linearly with the number of resource types. The RTFs will be described further in Section 5.

Penalty Functions

Penalty functions are designed to represent the user-level goals of the application. They are similar to priorities but are functions of the response time rather than a simple value so they can explicitly represent deadlines. Knowing the deadlines lets the system make optimizations that are difficult in today's systems such as running just fast enough to make the deadline. Like priorities, the penalty functions are set by the system on behalf of the user.

A ORADIA penalty function π is a non-decreasing piecewise linear function of the response time τ of the form

$$\pi(\tau) = \max(0, s(\tau - d)) \quad (2)$$

where d represents the deadline of the application and s (slope) defines the rate the penalty increases as response time increases. For applications without response time constraints the deadline can be set to 0. Two representative graphs of this type appear in Figures 2 and 3. Penalty Functions will be described further in Section 5.




Figure 2: A penalty function with a response time constraint.




Figure 3: A penalty function with no response time constraint.

Resource Allocation As Optimization

ORADIA formulates resource allocation as an optimization problem designed to minimize the total penalty of the system. This is analogous to minimizing user dissatisfaction with the user experience due to missed deadlines in a client system and minimizing the contract penalties paid for violated service-level agreements in a cloud system.

The optimization selects the allocations for all resources and resource types at once. This approach enables the system to make tradeoffs between resource types. For example, the system could choose to allocate more memory bandwidth in lieu of on-chip cache. It would difficult to provide predictable response times for applications considering the allocation of only one resource type at a time.

A succinct mathematical characterization of this resource allocation scheme is the following:

$$\begin{aligned} & \text{Minimize} && \sum_{p \in P} \pi_p(\tau_p(a_{p,1} \dots a_{p,n})) \\ & \text{Subject to} && \sum_{p \in P} a_{p,r} \leq A_r, r = 1, \dots, n \\ & && \text{and } a_{p,r} \geq 0 \end{aligned}$$

Here π_p is the penalty function for application p , τ_p is its response time function, $a_{p,r}$ is the allocation of resource r to application p , and A_r is the total amount of resource r available. Optimization details are described in Section 4.

Convex Optimization

If the penalty functions, response time functions, and resource constraints were arbitrary, little could be done to optimize the total penalty beyond searching at random for the best allocation. However, by framing our resource allocation problem as a *convex optimization problem* [9], two benefits accrue: an optimal solution will exist without multiple local extrema and fast, incremental solutions will become feasible.

A constrained optimization problem is convex if both the objective function to be minimized and the constraint functions that define its feasible solutions are convex functions. A function f is convex if its domain is a convex set and $F(\theta x + (1 - \theta)y) \leq \theta F(x) + (1 - \theta)F(y)$ for all θ between 0 and 1. A set is convex if for any two points x and y in the set, the point $\theta x + (1 - \theta)y$ is also in the set for all θ between 0 and 1. If F is differentiable, it is convex if its domain is an open convex set and $F(y) \geq F(x) + \nabla F^T \cdot (y - x)$ where ∇F is the gradient of F .

Therefore, a problem will be a convex optimization problem if it can be expressed in the form:

$$\begin{aligned} & \text{Minimize} && f_0(x_1, \dots, x_m) \\ & \text{Subject to} && f_i(x_1, \dots, x_m) \leq 0, i = 1, \dots, k \\ & && \text{where } \forall i \quad f_i : \mathfrak{R}^m \rightarrow \mathfrak{R} \text{ is convex.} \end{aligned}$$

As a consequence, ORADIA's resource allocation problem can be transformed into a convex optimization problem in the $m = |P| \cdot n$ variables $a_{p,r}$ as long as the penalty functions π_p are convex non-decreasing and the response time functions τ_p are convex. We designed our functions to meet these constraints, and proofs of their convexity can be found in Appendix B. Note that the resource constraints are all affine and can be rewritten as $\sum_{p \in P} a_{p,r} - A_r \leq 0$ and $-a_{p,r} \leq 0$.

ORADIA's convex formulation allows the time to solve the optimization to scale linearly with the number of resource types and the number of applications. For client operating systems with around 100 application running and 10 resource dimensions the total variables in the system to be solved is 1000—a very small problem which can be solved in microseconds on current systems. Cloud systems could possibly have

more applications running; however, the problem size scales linearly and the benefits gained could be significantly greater given the scale of the system.

Power and Energy

In ORADIA, we create an artificial application to represent the interests of system power and energy. Application 0 is designated the idle application and receives allocations of all resources that are left idle, *i.e.*, not allocated to other applications. The idle resources can be powered off or put to sleep if possible to save power.

The “response time” for application 0, τ_0 , is artificially defined to be the total system power consumption. The penalty function π_0 establishes a system tradeoff between power and performance that will determine which resources are allocated to applications to improve performance and which are left idle. The penalty function π_0 can be used to keep total system power below the parameter d_0 to the extent the penalties of other applications cannot overcome its penalty slope s_0 . Both s_0 and d_0 can be adjusted to reflect the current battery charge in mobile devices. For example as the battery depletes, s_0 could be increased to force other applications to slow or cease execution.

Additionally Application 0 functions as *slack* variables in our optimization problem turning the resource bounds into equalities:

$$\sum_{p \in P} a_{p,r} - A_r = 0, r = 1, \dots, n. \quad (3)$$

Resources

In our client system, resources are anything that the system can provide QoS on in hardware or software: specifically we use cpus, network bandwidth, cache ways, and memory pages in our system. However, other resources could easily be added as well assuming they have QoS enforcement mechanisms. The other scenarios would have resources that perform similar functions (compute, network, capacity), but at a different scale. For warehouse-scale computing resources are more likely to be different types of nodes, network bandwidth, and storage, and distributed embedded systems would include compute devices, link bandwidth, and memory.

3. Evaluation

In this section, we show the high-level results on how well ORADIA performs resource allocation to illustrate the viability of ORADIA before diving into all the system details in the following sections. Evaluations of the performance of individual components of ORADIA is included with the detailed technical descriptions of the components.

We use two different implementations of the ORADIA framework for evaluation. Our static framework collects data online using an x86 processor running Linux. That data is then used to build models and make and evaluate static resource allocation decisions. Our static framework was used

to experiment with the accuracy of different types of models and test the quality of the resource allocation decisions. Our dynamic framework is implemented in a research operating system, ROS, and is used to test our implementations of the algorithms, measure the overhead and reaction times, and illustrate ORADIA’s ability to work in a real system.

Static Resource Allocation Framework

Our static framework was designed to test the effectiveness of ORADIA’s model-based convex optimization for allocating resources. Data is collected online by running application benchmarks on a recent x86 processor running Linux-2.6.36. The measured data is processed using Python and then fed to MATLAB [46] to build the RTFs. MATLAB uses the RTFs to make resource allocation decisions. We compare performance of the chosen resource allocations with the actual measured performance of all possible resource allocations to test quality of the resource allocation decisions. We use CVX [17] in MATLAB to perform the convex optimization for building RTFs and making resource allocation decisions. We used this approach because it enabled us to test many applications, 44 in total, and many resource allocations rapidly.

Platform To collect data, we use a prototype version of Intel’s Sandy Bridge x86 processor that is similar to the commercially available client chip, but with additional hardware support for way-based LLC partitioning. The Sandy Bridge client chip has 4 quad-issue out-of-order superscalar cores, each of which supports 2 hyperthreads using simultaneous multithreading [32]. Each core has private 32 KB instruction and data caches, as well as a 256 KB private non-inclusive L2 cache. The LLC is a 12-way set-associative 6 MB inclusive L3 cache, shared among all cores using a ring-based interconnect. All three cache levels are write-back. The cache partitioning mechanism is way-based and works by modifying the cache-replacement algorithm. To allocate cache ways, we assign a subset of the 12 ways to a set of hyperthreads, thereby allowing only those hyperthreads to replace data in those ways.

We use a customized BIOS that enables the cache partitioning mechanism, and run unmodified Linux-2.6.36 for all of our experiments. To allocate cores, we use the Linux `taskset` command to pin applications to sets of hyperthreads. For our experiments we consider each hyperthread to be an independent core. To minimize the interference between hyperthreads, we first assign both hyperthreads available in one core before moving on to the next core. For example, allocations with four cores correspond to running on both hyperthreads of two cores.

Performance and Energy Measurement To measure application performance, we use the `libpfm` library [20, 53], built on top of the `perf_events` infrastructure introduced in Linux 2.6.31, to access available performance counters [33].

To measure on-chip energy, we use the energy counters available on Sandy Bridge to measure the consumption of the entire socket and also the total combined energy of cores, their

private caches, and the LLC. We access these counters using the Running Average Power Limit (RAPL) interfaces [33]. The counters measure power at a $1/2^{16}$ second granularity.

Description of Workloads Our workload contains a range of applications from three different popular benchmark suites: SPEC CPU 2006 [59], DaCapo [7], and PARSEC [6]. We selected this set of applications to represent a wide variety of possible resource behaviors in order to properly stress ORADIA’s RTFs. We include some additional applications to broaden the scope of the study, and some microbenchmarks to exercise certain system features.

The **SPEC CPU2006** benchmark suite [59] is a CPU-intensive, single-threaded benchmark suite, designed to stress a system’s processor, memory subsystem and compiler. Using the similarity analysis performed by Phansalkar et al. [54], we subset the suite, selecting 4 integer benchmarks (*astar*, *libquantum*, *mcf*, *omnetpp*) and 4 floating-point benchmarks (*cactusADM*, *calculix*, *lbm*, *povray*). Based on the characterization study by Jaleel [36], we also pick 4 extra floating-point benchmarks that stress the LLC: *GemsFDTD*, *leslie3d*, *soplex* and *sphinx3*. When multiple input sets and sizes are available, we pick the single *ref* input indicated by Phansalkar et al. [54].

We include the **DaCapo** Java benchmark suite as a representative of a managed-language workloads. We use the latest 2009 release, which consists of a set of open-source, real-world applications with non-trivial memory loads, and includes both client and server-side applications.

The **PARSEC** benchmark suite is intended to be representative of parallel real-world applications [6]. PARSEC programs use various parallelization approaches, including data- and task-parallelization. We use native input sets and the *pthread*s version for all benchmarks, with the exception of *freqmine*, which is only available in OpenMP.

We add four **additional parallel applications** to help ensure we cover the space of interest: *Browser_animation* is a multithreaded kernel representing a browser layout animation; *G500_csr* code is a breadth-first search algorithm; *Paradecoder* is a parallel speech-recognition application that takes audio waveforms of human speech and infers the most likely word sequence intended by the speaker; *Stencilprobe* simulates heat transfer in a fluid using a parallel stencil kernel over a regular grid [37].

We also add two **microbenchmarks** that stress the memory system: *stream_uncached* is a memory and on-chip bandwidth hog that continuously brings data from memory without caching it, while *ccbench* explores arrays of different sizes to determine the structure of the cache hierarchy.

Using a performance characterization of the applications, we select a subset of the benchmarks that are representative of different possible responses to resource allocations in order to reduce our study to a feasible size. Similar to [54], we use machine learning to select representative benchmarks. We use a hierarchical clustering algorithm [54] provided by the Python

library *scipy-cluster* with the *single-linkage* method. The feature vector contains parameters to represent core scaling, cache scaling, prefetcher sensitivity and bandwidth sensitivity. The clustering algorithm uses Euclidean distance between vectors to determine clusters.

The clustering results in 6 clusters representing the following (applications at the cluster center are listed in parenthesis):

- no scalability, high cache utility, (*429.mcf*)
- no scalability, low cache utility, (*459.gemsFDTD*)
- high scalability, low cache utility, (*ferret*)
- limited scalability, high cache utility, (*fop*)
- limited scalability, low cache utility, (*dedup*)
- limited scalability, low bandwidth sensitivity, (*batik*)

RTF Experiments To test the effectiveness of our RTFs in capturing real application behavior, we measured each of our 44 total benchmarks running alone on the machine for all possible resource allocations of cache and cores. Cores can be allocated from 1-8 and cache from 1-12 resulting in 96 possible allocations for each application. We use a genetic algorithm design of experiments [5] to select 10 and 20 of the collected allocations to build the RTFs. We also experimented with building RTFs with more data points but found that they provided little improvement over 20. We then use the model to predict the performance of every resource allocation and compare it with the actual measured performance of that resource allocation. We ran 3 complete trials and took the average error of the 3 trials.

Figure 4 shows the L1 norm of relative error of the predicted response times.

Resource Allocation Experiments Using the RTFs functions built for our

Figure 5 shows how ORADIA’s decisions compared with the optimal allocation, the Linux scheduler, and equally dividing the machine for each pair of the cluster representative applications. Figure 6 shows the same but using all pairs of PARSEC applications.

Dynamic Resource Allocation in a Manycore OS

In order to evaluate ORADIA’s effectiveness dynamically making decisions in real-time in a real operating system, we implemented in an in-house research operating system, ROS. We chose to implement in ROS over Linux for two reasons:

- ROS is designed around two-level scheduling meaning it more accurately reflects the system design ORADIA assumes
- ROS implements additional resource partitioning and QoS mechanisms allowing ORADIA to explore allocating more resource types

Platform Our dynamic experiments are all run on an Intel Nehalem-EP system with two 2.66-GHz Xeon X5550 quad-core processors and hyperthreading enable (*i.e.*, 16 hardware threads). It also contains a 1-Gbps Intel Pro/1000 Ethernet network adapter. ROS allocates resources directly to applications and then applications employ a second-level scheduler

Figure 4: L1 norm of relative error from RTF predicted response time compared to actual measured response time. Data represents the average error from 3 trials.

Figure 5: Resource allocation decisions for each pair of the cluster representative applications compared with the optimal allocation, the Linux scheduler, and equally dividing the machine.

Figure 6: Resource allocation decisions for each pair of PARSEC applications compared with the optimal allocation, the Linux scheduler, and equally dividing the machine.

to schedule threads onto the resources. There are two scheduling frameworks available in ROS, a cooperative framework called Lithe (LIquid THrEads) [52] and a preemptive one called PULSE (Preemptive User-Level SchEduling). All of our experiments are run using 2 different schedulers in the PULSE framework; applications with responsiveness requirements use the EDF, earliest-deadline first, scheduler and throughput oriented applications use the GRR, global round-robin, scheduler.

ROS in addition to allocating cores and cache ways as with our static framework, ROS can also allocate fractions of network bandwidth. In our experiments, we show ORADIA allocating cores and network bandwidth because the experimental platform, which has the advantage of more cores, does not have cache partitioning. We have also experimented with additional resources such as memory pages and cpu utilization, but they are not shown in this paper.

Performance and Energy Measurement Applications report their own measured response times to ORADIA through a message-passing interface built into ROS. ORADIA uses this information to build models, and we also use this information to show if the application is making its deadlines for the experiments.

ROS enables ORADIA to directly measure the system energy. However, energy counters are not available on our Nehalem-EP system and thus we extend the power model from the Sandy Bridge system to function as our Application 0 RTF.

Description of Workloads Our workload is designed to represent the scenario of a video conference, where each participant requires a separate, performance guaranteed video stream. New participants may join the conference and others leave, increasing or decreasing the number of streams running at any given time. Additionally, the speaker’s video is larger and higher resolution than the other video streams and as the speaker changes the requirements for video streams change. While conferencing, compute-intensive tasks, such as virus scans or file indexing, could be executed in background.

Our streaming video application is a multi-threaded, cpu- and network-intensive workload intended to simulate video chat applications like Skype and Facetime. Each video stream is encoded offline in the H.264 format using libx264, transported across the network through a TCP connection, and

decoded and displayed by the ROS client. The client receives, decodes and displays each frame using libffmepg and libx264. Each video stream has a corresponding EDF-scheduled thread with 33 ms deadlines using our second-level EDF scheduler.

We use psearchy [44], a parallel text indexer, from MOS-BENCH [10] as our file indexing application. We also use a network bandwidth hog application designed to contend with the video player for bandwidth by constantly sending UDP messages to the Linux box.

Resource Allocation Experiments Figure 7 shows the resource allocations for our video conference scenario when the computer is using wall power. Figure 8 shows same scenario except the computer is now using battery power and thus we have increased the importance of the wall.

4. Dynamic Penalty Optimization

ORADIA’s penalty optimization algorithm dynamically decides resource allocations. The algorithm can be run periodically, when applications start or stop, when an application changes faces or when the system changes operating scenarios. One of the advantages of convex optimization is that it enables fast, incremental solutions. As a result the algorithm could terminate earlier to decrease overhead and still be moving towards an optimal solution as it runs. However we found in our implementation that the algorithm was fast enough to run to completion every time.

Convex optimization is simplest when it is unconstrained, and so we reformulated ORADIA’s construction to be unconstrained. Extending the response time model functions to all of \mathcal{R}^n moves the requirement that allocations must be positive into the objective function, and introducing Application 0 for slack resources turns the affine inequalities into equalities:

$$\begin{aligned} &\text{Minimize} && \sum_{p \in P} \pi_p(\tau_p(a_{p,1} \dots a_{p,n})) \\ &\text{Subject to} && \sum_{p \in P} a_{p,r} = A_r, r = 1, \dots, n \end{aligned}$$

The only remaining constraints are those on the $a_{p,r}$. These can be removed by letting the $a_{p,r}$ be unbounded above for $p \neq 0$ and changing the domain of τ_0 to be the whole resource

Figure 7: Resource Allocations of video conferencing through time as participants enter and leave and the video in focus changes with the computer operating in wall power mode

Figure 8: Resource Allocations of video conferencing through time as participants enter and leave and the video in focus changes with the computer operating in battery power mode

allocation matrix. The definition of τ_0 might take the form

$$\begin{aligned}\tau_0 &= \sum_r \Delta_r \sum_{p \neq 0} a_{p,r} \\ &= \sum_r \Delta_r (A_r - a_{0,r})\end{aligned}$$

where Δ_r is the (constant) power dissipation of one unit of resource r . However, if any of the allocations $a_{0,r}$ turns out to be negative then τ_0 should instead return the value $+\infty$. This modification of the objective function transforms the resource allocation problem to unconstrained convex optimization.

The penalty optimization algorithm used in ORADIA is gradient descent via backtracking line search along the negative gradient direction [9]. This algorithm rejects and refines any step that yields insufficient relative improvement in the objective function, so infinite values from infeasible allocations will automatically be avoided by the search. The negative gradient $-\nabla\pi$ of the overall objective function π with respect to the resource allocations a is computed analytically from the response time models and penalty functions. When a component of this overall gradient is negative, it means the penalty will be reduced by increasing the associated allocation if possible. The gradient search at the boundaries of the feasible region must ignore components that lead in infeasible directions; these can be detected by noting whether for some p and r , $a_{p,r} = 0$ with $(-\nabla\pi)_{p,r} > 0$. In such cases, the associated step component is set to zero.

We added an additional optimization to move along boundaries more rapidly in the scenario when a completely allocated resource had a large gradient. We scale all the allocations of that resource type down to satisfy resource constraint while leaving the allocations of other resources untouched.

The rate of convergence of gradient descent depends on how well the sublevel sets of the objective function are conditioned (basically, how “spherical” they are). Conditioning will improve if resource allocation units are scaled to make their relative effects on t similar. For example, when compared with processor allocation units, memory allocation units of 4MB are probably a better choice than 4 KB. In addition, penalty function slopes should not differ by more than perhaps two orders of magnitude. If these measures prove insufficient, stronger preconditioners can be used. In our implementation, we condition all resource allocations to be in the range of 0-50.

Figure 9 shows performance of penalty optimization algorithm implemented in ROS. For our video conferencing scenario the average runtime is **x** and then worst case runtime is **y**. **explain where the variance comes from**

5. Application-Specific Functions

In this section, we describe the design of ORADIA’s response time and penalty functions in more detail.

Response Time Functions

A response time function describes an application’s performance based on its resource assignments. These functions capture information about the value of a particular resource type to an application. How well an application scales on a particular resource is high depend on the specific application and how it was written along with performance properties of the resources, which is why we believe it is necessary to collection application-specific data for each application. Without information of this type, it would be difficult for any resource allocation system informed decisions. Thus forcing the system to be entirely reactive and most likely require a complex set of heuristics.

Model Selection While it might have been possible to model response times by recording past values and interpolating among them, this idea has serious shortcomings:

- The size of the multidimensional response time function tables will be large;
- Interpolation in many dimensions is computationally expensive;
- The measurements will be noisy and require smoothing;
- Convexity in the resources may be violated;
- Gradient estimation will be slow and difficult.

Instead of interpolating, ORADIA maintains a parameterized analytic response time model with the partial derivatives evaluated from the model *a priori*. Application responsiveness is highly nonlinear for an increasing variety of applications like streaming media or gaming, thus requiring many data points to represent the response times without a model. Using models each application can be described in a small number of parameters. Additionally models can be built from few data points and naturally smooth out noisy data. The resource gradient can be easily calculated which is essential for ORADIA to solve the optimization problem efficiently.

ORADIA models response times by functions that are convex by construction. The specific function chosen for ORADIA is shown in Equation 1. In this equation, the response time is modeled as a weighted sum of component terms, one per resource, where each term w_r/a_r is the amount of work $w_r \geq 0$ divided by a_r , the allocation of that resource [58]. For example, one term might model the number of instructions executed divided by total processor MIPS, another might model

Figure 9: Performance of our penalty optimization algorithm

storage accesses divided by storage allocation and so forth. Asynchrony and latency tolerance may make response time components overlap partly or fully; and thus we added additional terms to represent the interactions between resources. We choose this specific function because in initial application studies we found it models response time behavior accurately enough to allow the optimization to make good decisions but is also low overhead to build. Alternative models and the initial model evaluations are described in Appendix A.

Such models are automatically convex in the allocations because $1/a$ is convex for positive a and because a positively-weighted sum of convex functions is convex. It is obviously important to guarantee the positivity of the resource allocations. This can be enforced as the allocations are selected during penalty optimization, or the response time model can be made to return ∞ if any allocation is less than or equal to zero. This latter idea preserves the convexity of the model and extends its domain to all of \mathbb{R}^n and consequently we used this approach in our implementation. The gradient $\nabla \tau$ is needed by the penalty optimization algorithm. Since τ is analytic, generic, and symbolically differentiable it is a simple matter to compute the gradient of τ once the model is defined.

Convexity Assumption Forcing our response time functions to be convex assumes that the actual response time are reasonably convex; We find this is a plausible requirement as applications almost completely follow the proverbial “Law of Diminishing Returns” for resource allocations.

However, there are examples of response time versus resource behavior that violate convexity. For example, we have particularly seen non-monotonic performance behavior in applications when dealing with hyperthreads or memory pages. For two of our applications 5 hyperthreads resulted in significantly worse performance than 4 or 6 hyperthreads. When studying some other applications we found that particular numbers of memory pages, *i.e.*, 2K, resulted in much better performance than the adjacent page allocations. One possible solution to solve this problem is keep track of points with significant error in the models and adjust resource allocations slightly to avoid these points.

Another such example sometimes occurs in memory allocation, where “plateaus” can sometimes be seen as in Figure 10. Such plateaus are typically caused by algorithm adaptations within the application to accommodate variable resource availability. The response time is really the *minimum* of several convex functions depending on allocation and the point-wise minimum that the application implements fails to preserve convexity. The effect of the plateaus will be a non-convex penalty as shown in Figure 11 and multiple extrema in the optimization problem will be a likely result. There are several ways to avoid this problem. One is based on the observation that such response time functions will at least be *quasiconvex*.

Another idea is to use additional constraints to explore convex sub-domains of τ .




Figure 10: Response time function with some resource “plateaus”.




Figure 11: Net effect of the resource plateaus on the application penalty.

Forcing our response time functions to be convex assumes that the actual response time are reasonably convex; We find this is a plausible requirement as applications almost completely follow the proverbial “Law of Diminishing Returns” for resource allocations.

However, there are examples of response time versus resource behavior that violate convexity. For example, we have particularly seen non-monotonic performance behavior in applications when dealing with hyperthreads or memory pages. For two of our applications 5 hyperthreads resulted in significantly worse performance than 4 or 6 hyperthreads. When studying some other applications we found that particular numbers of memory pages, *i.e.*, 2K, resulted in much better performance than the adjacent page allocations. One possible solution to solve this problem is keep track of points with significant error in the models and adjust resource allocations slightly to avoid these points.

Another such example sometimes occurs in memory allocation, where “plateaus” can sometimes be seen as in Figure 10. Such plateaus are typically caused by algorithm adaptations within the application to accommodate variable resource availability. The response time is really the *minimum* of several convex functions depending on allocation and the point-wise minimum that the application implements fails to preserve convexity. The effect of the plateaus will be a non-convex penalty as shown in Figure 11 and multiple extrema in the optimization problem will be a likely result. There are several ways to avoid this problem. One is based on the observation that such response time functions will at least be *quasiconvex*. Another idea is to use additional constraints to explore convex sub-domains of τ .

Overall, we found that our simple convex models still resulted in high-quality resource allocations and thus chose not to implement any of these approaches. Alternative approaches to handling non-convex behavior are described in Appendix B. Additional challenges to response time modeling are discussed in Section 7.

Data Collection and Creation Time There are several possible ways to collect the response time data for applications. A user-level runtime scheduler that schedules work internal to the application may be a good source of data or the operating system could measure progress using performance counters. In our implementation applications report their own measured values: however, this solution was chosen simply as a way to test the validity of the concept. In a production operating system, it may not be the best approach because it requires trusting applications not to lie about their performance. In a datacenter environment this may be less of a concern.

There are also many different possible moments response time functions could be created. RTFs could be created in advance and distributed with the application. In the case of app stores this approach could make lot of sense since most app stores only cater to a limited number of platforms. Data could also be crowd sourced and the RTFs built in the cloud, which has the advantage making it easy to collect a diverse set of training points. However, all of these approaches are missing personalization. As a result we have chosen to implement two solutions that collect data directly from the users machine. The first approach is to collect all of the training points at application install time and build the model then. The more advanced approach collects data continuously as the application runs and adds that data to the model training set and rebuilds the model. Although ultimately a hybrid approach may be the most effective: applications can begin with a provided generic model, and the system can improve the model over time. Section 6 describes our model creation process in detail.

Penalty Functions

Penalty functions are generically defined as members of a family of such functions so that user preferences for an application p can be implemented by assigning values to a few well-understood parameters. We designed ORADIA's penalty functions to represent deadlines and the relative importance of making them. These values are represented by two parameters a deadline d and a slope s . Equation 2 shows penalty functions used in ORADIA.

A response-time constrained application has a marked change in slope, namely from 0 to s , at the point $\tau = d$. In the most extreme case $s = \infty$ (implying infinite penalty for the system as a whole when $\tau > d$). For applications without response time constraints one can set $d = 0$. This defines linear behavior with s as the rate of penalty increase with response time.

To guarantee the penalty is convex and non-decreasing, s must be non-negative. The response time τ is of course non-negative, and it may be sensible (if not strictly necessary) to require that d is also.

The gradient of application penalty with respect to its resource allocations is used in the optimization algorithm. By the chain rule, $\partial\pi/\partial a_r = \partial\pi/\partial\tau \cdot \partial\tau/\partial a_r$. The first term

is well-defined albeit discontinuous at $\tau = d$ with $\partial\pi/\partial\tau =$ if $(\tau - d) \leq 0$ then 0 else s .

In a client operating system, the instantaneous management of penalty function modifications should be highly automated by the system to avoid unduly burdening the user. Most easily parameters could be set similarly to how priorities are set in some systems today. Applications are grouped into *interaction classes* based on application type and the interaction class defines the deadline and slope. In our implementation we set the parameters manually. However, future work is to experiment with alternative methods for *learning* the deadline and slope.

As a application grows or diminishes in importance, its penalty function can be modified accordingly. Penalty function adjustment is most likely to occur in transitions between operating scenarios. For example, when unplugging a device all of the background activities could have their slopes significantly reduced to save battery life.

Power and Battery Energy

In order to control power and energy of the system, we use an artificial application named Application 0 which receives all resources not allocated to other applications. Application 0's response time function is similar to the other applications' response time functions. The function inputs are resource allocations just as with the other applications. However, the function output is system power rather than response time. To create the RTF, system power can be measured directly from on-chip energy counters in systems where they are available or from a power meter. These models can be built in advance, during a training phase or online while the system runs, just as with the application RTFs. Alternatively, the model could be part of the operating system platform-specific information.

Although system power may not be entirely convex in reality, approximating it to be convex is reasonable because the convex model still captures the general behavior that leaving a resource idle should save some or no power. As a result, Application 0 still fills its purpose of preventing applications using additional resources that have low performance/power ratios.

The penalty function can be adjusted to represent how important saving power and energy is in the current operating scenario. For example, if a mobile device is plugged in to a charger than perhaps the penalty slope should be set very low. However if the device is unplugged the penalty function could be changed to have a deadline and then an steep slope which represents creating a power cap for the system. In this way the deadline and slope could be set to help control the minimum number of hours the battery will last.

6. Response Time Model Creation

To create RTF models either at install time or online we use a least squares approach described below. For the install time approach we use a genetic algorithm, Audze-Eglasis Design of Experiments [5], to select the resource vectors to use for

training data. The data for all those vectors is collected and fed into the least squares algorithm. In the online approach the data comes from the application's response time history.

Model Creation using Least Squares After sufficient measurements, discovery of the model parameters w that define the function τ can be based on a solution to the over-determined linear system $t = Dw$ where t is a column vector of actual response times measured for the application and D is a matrix whose i th row $D_{i,*}$ contains the reciprocals of the amplified bandwidth allocations that generated the corresponding response time measurement t_i . Estimating w is relatively straightforward: a least-squares solution accomplished via *QR factorization* [24] of D will determine the w that minimizes (half the) square of the residual error $1/2\|Dw - t\|_2^2 = 1/2\|\varepsilon\|_2^2$. The solution proceeds as follows:

$$\begin{aligned} t &= Dw - \varepsilon \\ &= QRw - \varepsilon \\ Q^T t &= R w - Q^T \varepsilon \end{aligned}$$

It is not always necessary to materialize the orthogonal matrix $Q^T = Q^{-1}$; the individual elementary orthogonal transformations (Householder reflections or Givens rotations) that triangularize R by progressively zeroing out partial columns of D can simultaneously be applied to t . The elements of the resulting vector $Q^T t$ that correspond to zero rows in R comprise $-Q^T \varepsilon$. Since Rw exactly equals the upper part of $Q^T t$, the upper part of $Q^T \varepsilon$ is zero. The residual error for the t_i can be found by premultiplying $Q^T \varepsilon$ by Q .

This formulation assumes a model norm $p = 1$. If a different model norm p is desired, such as $p = 2$, we could first square each measurement in t and each reciprocal bandwidth term in D and then follow the foregoing procedure. The elements of the result w will be squares as well, and the 2-norm of the difference in the squared quantities will be minimized. This is not the same as minimizing the 4-norm; what is being minimized is $1/2\|\text{diag}(Dww^T D^T - tt^T)\|_2^2$.

On-line Response Time Modeling

As resource allocation continues, more measurements will become available to augment t and D . Moreover, older data may become a poor representation of the current behavior of the application if its characteristics have changed, presumably as reflected in $Q^T \varepsilon$. ORADIA adopts the incremental least squares described below to replace old data and efficiently update RTFs.

Incremental Least Squares What is needed is a factorization $\tilde{Q}\tilde{R}$ of a new matrix \tilde{D} derived from D by dropping a row, perhaps from the bottom, and adding a row, perhaps at the top. Corresponding elements of t are dropped and added to form \tilde{t} .

The matrices \tilde{Q} and \tilde{R} can be generated by applying Givens rotations in the way described in Section 12 of [24] to *down-date* or *update* the factorization much more cheaply than re-computing it *ab initio*. The method requires retention and

maintenance of Q^T but not of D . Every update in ORADIA is preceded by a downdate that makes room for it. Dated rows are *not* always the oldest (bottom) ones, but an update always adds a new top row. For several reasons, the number of rows m in R will be at least twice the number of columns n . Rows selected for downdating will always be in the lower $m - n$ rows of R , guaranteeing that the most recent n updates are always part of the model.

Non-Negativity To guarantee convexity of the response time model, the solution w to $t \approx QRw$ must have no negative components. Intuitively, when a resource is associated with more than a single w_j or when the measured response time increases with allocation then negative w_j may occur.

A requirement for non-negative solutions to least-squares linear algebra problems is common, so much so that it has a name: *Non Negative Least Squares*, or NNLS. There are several well-known techniques [13], but since the method proposed here for online model maintenance calls for incremental downdates and updates to rows of Q^T , $Q^T t$ and R , the NNLS problem is handled with a complementary scheme based on the *active-set* method [42] that also downdates and updates the *columns* of R incrementally, roughly in the spirit of Algorithm 3 in [45]. However, ORADIA's algorithm cannot ignore downdated columns of R because subsequent *row* updates and downdates must have due effect on these columns to allow their later reintroduction via column updates when necessary. /pacora solves this problem by leaving the downdated columns in place, skipping over them in maintaining and using the QR factorization.

Figure 12 shows the performance of ORADIA's NNLS algorithm.

Model Rank Preservation If care is not taken in the allocation process, the rows of R may become linearly dependent to such an extent that its rank is insufficient to determine w . This might be the result of repetitions in resource assignment updates. There are several possible ways to avoid this *rank-deficiency* problem. The characteristics of R depend on both the resource optimization trajectory and the choices made in the downdate-update algorithm. In particular, deciding whether to downdate the bottom row of R or some "younger" row will depend on whether the result would become rank-deficient. This approach decouples allocation optimization from performance model maintenance and places responsibility upon the latter to always keep enough history to determine a model.

Outliers and Phase Changes Some response time measurements may be "noisy" or even erroneous. A weakness of least-squares modeling is the high importance it gives to outlying values. On the other hand, when an application changes phase it is important to adapt quickly, and what looks like an outlier when it first appears may be a harbinger of change. What is needed is a way to discard either old or outlying data with a judicious balance between age and anomaly.

The downdating algorithm accomplishes this by weighting

Figure 12: Performance of ORADIA's NNLS algorithm

Figure 13: ORADIA's online model creation algorithm adapting to a video application changing phases

the errors in $\varepsilon = Q(Q^T t - R w)$ between the predicted response times τ and the measured ones t by a factor that increases exponentially with the age $g(i)$ of the error ε_i . Age can be modeled coarsely by the number of time quanta of some size since the measurement; ORADIA simply lets $g(i) = i$. The weighting factor for the i th row is then $\eta^{g(i)}$ where η is a constant somewhat greater than 1. The candidate row to downdate is the row with the largest weighted error, *i.e.*

$$dd = \arg \max_i |\varepsilon_i| \cdot \eta^{g(i)}$$

Figure 13 shows ORADIA's model creation algorithm adapting to a video application changing phases.

7. Discussion

ORADIA for Warehouse-Scale Computing

Discuss major differences, particularly storage

Challenges

Outliers and Performance Non-Monotonicity Figure 14 shows the effect of model accuracy on the quality of the resource allocation decisions made.

Discuss Heuristics + Feedback

Variability Discuss phases, input dependence, other variability (such as network connection), shared resources

Possible solutions include Heuristics + Feedback, Stochastic Models, Changing Models

8. Related Work

Add Gang Scheduling? Discuss Application Performance Modeling

8.1. Autonomic Systems

In much of this research area [39, 49, 62, 63] the performance models are derived from off-line measurements and are either interpolations from tables or analytic functions based on queueing theory. The utility functions typically map the number of servers each execution environment receives to its performance relative to its requirements, which may be multiple. A central arbiter maximizes total utility. The utility functions are not necessarily concave, so the arbiter must use reinforcement learning or combinatorial search to make allocations. Walsh *et al.* [64] note the importance of basing utility functions on the metrics in which QoS is expressed rather than on the raw quantities of resources. Each application has a manager that schedules the resources given to it by a global arbiter. There are other philosophical similarities to ORADIA,

but since the objective functions are neither continuous nor convex their optimization is difficult.

The related systems AcOS [4] and Metronome [57] feature hardware-thread based maintenance of “Heart Rate” targets using adaptive reinforcement learning techniques. AcOS also senses thermal conditions and can exploit Dynamic Voltage and Frequency Scaling (DVFS).

A survey of autonomic systems research appears in [30].

8.2. Soft Real-Time and SLAs

Soft real-time systems typically modify an existing operating system to improve responsiveness. Calandrino *et al.* [11] use working set sizes to make co-scheduling decisions and enhance soft real-time behavior.

Rajkumar *et al.* [56] propose a system Q-RAM that maximizes the weighted sum of utility functions, each of which is a function of the resource allocation to the associated application. Unlike ORADIA, there is no distinction between performance and utility, and the utility functions are assumed as input rather than being discovered by the system. They are always nondecreasing in every resource and sometimes are piecewise-linear and concave; in these cases the optimal allocation is easily found by a form of gradient ascent. When the utility functions are not concave, a suboptimal greedy algorithm is proposed.

In the Redline system [66], compact resource requirement specifications are written by hand to guarantee response times. Isolation of resources is strong, as in ORADIA. Scheduling is Earliest-Deadline-First. Admission control is lenient but oversubscription situations are remedied by de-admitting some of the non-interactive applications.

Jockey [21] has some similarities to ORADIA: it is intended to handle parallel computation, its utility functions are concave, and it adapts dynamically to application behavior. Its performance models are obtained by calibrating either event-based simulation or a version of Amdahl's Law to computations. Jockey does not optimize total utility but simply increases processors until utility flattens for each application, *i.e.* each deadline is met. A fairly sophisticated control loop prevents oscillatory behavior.

Bodik *et al.* [8] is also like ORADIA in that it builds online performance models. Initially, it uses an *exploration policy* that avoids nearly all SLA violations while it builds the model; later, it shifts to a controller based on the model it has built. The models are statistical, and bootstrapping is used to estimate performance variance. Major changes in the application model are detected and cause model exploration to resume. The models are not convex or concave in general,

Figure 14: Effect of Model Accuracy on Decision Quality

and all SLAs must be met with high probability.

8.3. Partitioning Mechanisms

Most hardware partitioning mechanism work has looked at shared cache structures and provided mechanisms to partition them according to a varied set of goals. Suh *et al.* [60, 61] and Qureshi and Patt [55] monitor individual applications' cache performance and use this monitoring to inform their partitioning mechanism in an attempt to reduce the total amount of cache misses and off-chip memory traffic. A wide variety of proposals exist for multicore last-level cache structures that partition the spatial resources between private and shared data, in an attempt to create a manageable trade-off between capacity for shared data and low latency for private data [12, 15, 19, 31, 47, 68, 69].

Further cache partitioning work has focused on providing QoS guarantees to applications. Early work focused on providing adaptive, fair policies that ensure equal performance degradation [38, 67], while more recent proposals have incorporated more sophisticated policy management [25, 26, 29, 35]. Other partitioning work has focused on interconnect bandwidth QoS [43] or partitioning cache capacity and bandwidth simultaneously [50]. In general, these papers focus on designing and proving the effectiveness of particular mechanisms for particular goals, without a concrete notion of a general framework in which a variety of application-specific QoS requirements can be communicated to an all-purpose resource allocator and scheduler. For example, Iyer [34] suggests a priority-based cache QoS framework, CQoS, for shared cache way-partitioning. The priorities might be specified per core, per application, per memory type, or even per memory reference. Simultaneous achievement of performance targets is not addressed.

8.4. Resource Allocation Frameworks

Guo *et al.* [26] point out that much prior work is insufficient for true QoS – merely partitioning hardware is not enough, because there must also be a way to specify performance targets and an admission control policy for jobs. The framework they present incorporates a scheduler that supports multiple execution modes.

Nesbit *et al.* [51] introduce the *Virtual Private Machines* (VPM), a framework for resource allocation and management in multicore systems. A VPM consists of a set of virtual hardware resources, both spatial (physical allocations) and temporal (scheduled time-slices). Unlike traditional virtual machines that only virtualize resource functionality, VPMs virtualize a system's performance and power characteristics, meaning that a VPM has the same performance and power profile as a real machine with an equivalent set of hardware resources.

They break down the framework components into policies and mechanisms which may be implemented in hardware or software. Critical software components of the VPM framework are VPM *modeling*, which maps high-level application objectives into VPM configurations, and VPM *translation*, which uses VPM models to assign an acceptable VPM to the application while adhering to system-level policies. A VPM scheduler then decides if the system can accommodate all applications or whether resources need to be revoked.

The VPM approach and terminology mesh well with our study, which can be seen as a specific implementation of several key aspects of the type of framework they describe (i.e. VPM modeling and translation). Nesbit *et al.* did not perform any evaluations of the modeling, translation, or scheduling processes suggested in their paper.

The Singularity operating system [1] provides process isolation through software rather than hardware. This isolation in accessibility is not the same as the performance isolation that is desirable for ORADIA.

Kumar *et al.* [40] demonstrate the performance advantages of heterogeneous cores for mixed workloads using heuristic allocation strategies in both space and time.

Genbrugge and Eeckhout [23] demonstrate the importance of adapting to changes in application characteristics, in this case instructions per cycle.

Merkel and Bellosa [48] propose *Task Activity Vectors* that describe how much each application uses the various functional units; these vectors are used to balance usage across multiple cores and unbalance usage among hardware threads within each core. The intended effect is to distribute chip temperature more evenly, but the idea may be more broadly applicable, e.g. for heterogeneous systems.

Ganapathi *et al.* have had success using machine learning to model application performance and select the best performing configuration in [22].

[TODO: I DON'T KNOW WHY [27, 65] ARE IN HERE]

Chen *et al.* [14] propose parallel depth-first thread scheduling as an alternative to work-stealing for constructive cache sharing. The paper discusses the automation of task granularity selection to match cache capacity. This kind of tuning is particularly appropriate for a user-level runtime system.

9. Conclusion

Discuss advantages of ORADIA Possible advantages to discuss:

- Heterogeneity,
- Convex optimization is relatively inexpensive optimization problem with good techniques to solve,
- Penalty Function Slopes allow the system to express relative priorities of application.
- Priorities change as a function of performance (criticality),

- Penalty Function Intercept encapsulates QoS requirements (responsiveness),
- Additional process can be used to represent system energy,
- Works the growing diversity of modern hardware,
- doesn't require application developers to know about resources

References

- [1] M. Aiken *et al.*, "Deconstructing process isolation," in *Memory Systems Performance and Correctness*, 2006.
- [2] Apple Inc., "iOS App Programming Guide," <http://developer.apple.com/library/ios/DOCUMENTATION/iPhone/Conceptual/iPhoneOSProgrammingGuide/iPhoneAppProgrammingGuide.pdf>.
- [3] L. A. Barroso and U. Hözl, *The Datacenter as a Computer: An Introduction to the Design of Warehouse-Scale Machines*, ser. Synthesis Lectures on Computer Architecture. Morgan & Claypool Publishers, 2009.
- [4] D. B. Bartolini *et al.*, "Acos: an autonomic management layer enhancing commodity operating systems," in *In DAC Workshop on Computing in Heterogeneous, Autonomous, 'N' Goal-oriented Environments (CHANGE), co-located with the Annual Design Automation Conference (DAC)*, 2012.
- [5] S. J. Bates, J. Sienz, and D. S. Langley, "Formulation of the audze-eglais uniform latin hypercube design of experiments," *Adv. Eng. Softw.*, vol. 34, no. 8, pp. 493–506, 2003.
- [6] C. Bienia *et al.*, "The parsec benchmark suite: Characterization and the parsec benchmark suite: Characterization and architectural implications," Princeton University, Tech. Rep. TR-811-08, January 2008.
- [7] S. M. Blackburn *et al.*, "The DaCapo benchmarks: Java benchmarking development and analysis," in *OOPSLA*, 2006, pp. 169–190.
- [8] P. Bodik *et al.*, "Automatic exploration of datacenter performance regimes," in *Proc. ACDC*, 2009.
- [9] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge, England: Cambridge University Press, 2004.
- [10] S. Boyd-Wickizer *et al.*, "Mosbench," 2012, <http://pdos.csail.mit.edu/mosbench/>.
- [11] J. M. Calandrino and J. H. Anderson, "On the design and implementation of a cache-aware multicore real-time scheduler," in *ECRTS*, 2009, pp. 194–204.
- [12] J. Chang and G. S. Sohi, "Cooperative cache partitioning for chip multiprocessors," in *Proc. ICS '07*. New York, NY, USA: ACM, 2007, pp. 242–252.
- [13] D. Chen and R. J. Plemmons, "Nonnegativity constraints in numerical analysis," Lecture presented at the symposium to celebrate the 60th birthday of numerical analysis, Leuven, Belgium, October 2007.
- [14] S. Chen *et al.*, "Scheduling threads for constructive cache sharing on CMPs," in *Proc. SPAA '07*. New York, NY, USA: ACM, 2007, pp. 105–115.
- [15] S. Cho and L. Jin, "Managing distributed, shared L2 caches through os-level page allocation," in *Proc. MICRO 39*. Washington, DC, USA: IEEE Computer Society, 2006, pp. 455–468.
- [16] G. Contreras and M. Martonosi, "Characterizing and improving the performance of intel threading building blocks," in *Workload Characterization, 2008. IISWC 2008. IEEE International Symposium on*, Sept., pp. 57–66.
- [17] I. CVX Research, "CVX: Matlab software for disciplined convex programming, version 2.0 beta," <http://cvxr.com/cvx>, Sep. 2012.
- [18] J. Dean and S. Ghemawat, "Mapreduce: simplified data processing on large clusters," *Commun. ACM*, vol. 51, no. 1, pp. 107–113, Jan. 2008. Available: <http://doi.acm.org/10.1145/1327452.1327492>
- [19] H. Dybdahl and P. Stenstrom, "An adaptive shared/private nuca cache partitioning scheme for chip multiprocessors," in *Proc. HPCA '07*. Washington, DC, USA: IEEE Computer Society, 2007, pp. 2–12.
- [20] S. Eranian, "Perfmon2: a flexible performance monitoring interface for linux," in *Ottawa Linux Symposium*, 2006, p. 269–288.
- [21] A. D. Ferguson *et al.*, "Jockey: guaranteed job latency in data parallel clusters," in *Proceedings of the 7th ACM european conference on Computer Systems*, ser. EuroSys '12. New York, NY, USA: ACM, 2012, pp. 99–112. Available: <http://doi.acm.org/10.1145/2168836.2168847>
- [22] A. Ganapathi *et al.*, "A case for machine learning to optimize multicore performance," in *HotPar09*, Berkeley, CA, 3/2009 2009. Available: <http://www.usenix.org/event/hotpar09/tech/>
- [23] D. Genbrugge and L. Eeckhout, "Statistical simulation of chip multi-processors running multi-program workloads," in *ISCA*, 2007.
- [24] G. H. Golub and C. F. Van Loan, *Matrix Computations*, 3rd ed. Baltimore, Maryland: Johns Hopkins University Press, 1996.
- [25] F. Guo *et al.*, "From chaos to qos: case studies in cmp resource management," *SIGARCH Comput. Archit. News*, vol. 35, no. 1, pp. 21–30, 2007.
- [26] F. Guo *et al.*, "A framework for providing quality of service in chip multi-processors," in *Proc. MICRO '07*. Washington, DC, USA: IEEE Computer Society, 2007, pp. 343–355.
- [27] J. L. Hellerstein *et al.*, *Feedback Control of Computing Systems*. John Wiley & Sons, 2004.
- [28] J. L. Hennessy and D. A. Patterson, *Computer Architecture - A Quantitative Approach (5. ed.)*. Morgan Kaufmann, 2012.
- [29] L. R. Hsu *et al.*, "Communist, utilitarian, and capitalist cache policies on cmps: caches as a shared resource," in *Proc. PACT '06*. New York, NY, USA: ACM, 2006, pp. 13–22.
- [30] M. C. Huebscher and J. A. McCann, "A survey of autonomic computing—degrees, models, and applications," *ACM Comput. Surv.*, vol. 40, no. 3, pp. 1–28, 2008.
- [31] J. Huh *et al.*, "A nuca substrate for flexible cmp cache sharing," in *Proc. ICS '05*. New York, NY, USA: ACM, 2005, pp. 31–40.
- [32] Intel Corp., "Intel 64 and ia-32 architectures optimization reference manual," June 2011.
- [33] Intel Corp., "Intel 64 and ia-32 architectures software developer's manual," March 2012.
- [34] R. Iyer, "Cqos: a framework for enabling qos in shared caches of cmp platforms," in *Proc. ICS '04*. New York, NY, USA: ACM, 2004, pp. 257–266.
- [35] R. Iyer *et al.*, "Qos policies and architecture for cache/memory in cmp platforms," in *Proc. SIGMETRICS '07*. New York, NY, USA: ACM, 2007, pp. 25–36.
- [36] A. Jaleel, "Memory characterization of workloads using instrumentation-driven simulation – a pin-based memory characterization of the spec cpu2000 and spec cpu2006 benchmark suites," VSSAD, Intel Corporation, Tech. Rep., 2007.
- [37] S. Kamil, "Stencil probe," 2012, <http://www.cs.berkeley.edu/~skamil/projects/stencilprobe/>.
- [38] C. Kim, D. Burger, and S. W. Keckler, "An adaptive, non-uniform cache structure for wire-delay dominated on-chip caches," in *Proc. ASPLOS-X*. New York, NY, USA: ACM, 2002, pp. 211–222.
- [39] S. Kounev, R. Nou, and J. Torres, "Autonomic qos-aware resource management in grid computing using online performance models," in *Proc. ValueTools '07*. ICST, 2007, pp. 1–10.
- [40] R. Kumar *et al.*, "Single-isa heterogeneous multi-core architectures for multithreaded workload performance," in *ISCA '04: Proceedings of the 31st annual international symposium on Computer architecture*. Washington, DC, USA: IEEE Computer Society, 2004, p. 64.
- [41] H. T. Kung, "Memory requirements for balanced computer architectures," in *International Symposium on Computer Architecture*, 1986, pp. 49–54.
- [42] C. L. Lawson and R. J. Hanson, *Solving Least Squares Problems*. Englewood Cliffs, NJ: Prentice Hall, 1974.
- [43] J. W. Lee, M. C. Ng, and K. Asanovic, "Globally-synchronized frames for guaranteed quality-of-service in on-chip networks," in *Proc. ISCA '08*. Washington, DC, USA: IEEE Computer Society, 2008, pp. 89–100.
- [44] J. Li *et al.*, "On the feasibility of peer-to-peer web indexing and search," in *Peer-to-Peer Systems II*, ser. Lecture Notes in Computer Science, M. Kaashoek and I. Stoica, Eds. Springer Berlin Heidelberg, 2003, vol. 2735, pp. 207–215. Available: http://dx.doi.org/10.1007/978-3-540-45172-3_19
- [45] Y. Luo and R. Duraiswami, "Efficient parallel nonnegative least squares on multicore architectures," *SIAM Journal on Scientific Computing*, vol. 33, no. 5, pp. 2848–2863, 2011.
- [46] Mathworks, "Matlab 2009b," <http://www.mathworks.com/>, 2009.
- [47] J. Merino *et al.*, "Sp-nuca: a cost effective dynamic non-uniform cache architecture," *SIGARCH Comput. Archit. News*, vol. 36, no. 2, pp. 64–71, 2008.
- [48] A. Merkel and F. Bellosa, "Task activity vectors: a new metric for temperature-aware scheduling," in *Proc. Eurosys '08*. New York, NY, USA: ACM, 2008, pp. 1–12. Available: http://portal.acm.org/ft_gateway.cfm?id=1352594&type=pdf&coll=GUIDE&dl=GUIDE&CFID=35811817&CFTOKEN=17839388

- [49] M. N. Bannani and D. A. Menasce, "Resource allocation for autonomic data centers using analytic performance models," in *ICAC '05*. Washington, DC, USA: IEEE Computer Society, 2005, pp. 229–240.
- [50] K. J. Nesbit, J. Laudon, and J. E. Smith, "Virtual private caches," in *Proc. ISCA '07*. New York, NY, USA: ACM, 2007, pp. 57–68.
- [51] K. J. Nesbit *et al.*, "Multicore resource management," *IEEE Micro*, vol. 28, no. 3, pp. 6–16, 2008.
- [52] H. Pan, B. Hindman, and K. Asanović, "Lithe: Enabling efficient composition of parallel libraries," in *HotPar09*, Berkeley, CA, 03/2009 2009. Available: <http://www.usenix.org/event/hotpar09/tech/>
- [53] "Perfmon2 webpage," perfmon2.sourceforge.net/.
- [54] A. Phansalkar, A. Joshi, and L. K. John, "Analysis of redundancy and application balance in the spec cpu2006 benchmark suite," in *ISCA*, 2007, pp. 412–423.
- [55] M. K. Qureshi and Y. N. Patt, "Utility-based cache partitioning: A low-overhead, high-performance, runtime mechanism to partition shared caches," in *Proc. MICRO 39*. Washington, DC, USA: IEEE Computer Society, 2006, pp. 423–432.
- [56] R. Rajkumar *et al.*, "A resource allocation model for qos management," in *Proc. RTSS '97*. Washington, DC, USA: IEEE Computer Society, 1997, p. 298.
- [57] F. Sironi *et al.*, "Metronome: operating system level performance management via self-adaptive computing," in *Proceedings of the 49th Annual Design Automation Conference*, ser. DAC '12. New York, NY, USA: ACM, 2012, pp. 856–865. Available: <http://doi.acm.org/10.1145/2228360.2228514>
- [58] A. Snaveley *et al.*, "A framework for performance modeling and prediction," in *SC*, 2002, pp. 1–17.
- [59] Standard Performance Evaluation Corporation, "SPEC CPU 2006 benchmark suite," <http://www.spec.org>.
- [60] G. E. Suh, L. Rudolph, and S. Devadas, "Dynamic partitioning of shared cache memory," *J. Supercomput.*, vol. 28, no. 1, pp. 7–26, 2004.
- [61] G. E. Suh, S. Devadas, and L. Rudolph, "A new memory monitoring scheme for memory-aware scheduling and partitioning," in *Proc. HPCA '02*. Washington, DC, USA: IEEE Computer Society, 2002, p. 117.
- [62] G. Tesaro, W. E. Walsh, and J. O. Kephart, "Utility-function-driven resource allocation in autonomic systems," in *Proc. ICAC '05*. Washington, DC, USA: IEEE Computer Society, 2005, pp. 342–343.
- [63] G. Tesaro *et al.*, "On the use of hybrid reinforcement learning for autonomic resource allocation," *Cluster Computing*, vol. 10, no. 3, pp. 287–299, 2007.
- [64] G. Tesaro and J. O. Kephart, "Utility functions in autonomic systems," in *Proc. ICAC '04*. Washington, DC, USA: IEEE Computer Society, 2004, pp. 70–77.
- [65] L. Wasserman, *All of Nonparametric Statistics (Springer Texts in Statistics)*. Secaucus, NJ, USA: Springer-Verlag New York, Inc., 2006.
- [66] T. Yang *et al.*, "Redline: first class support for interactivity in commodity operating systems," in *Proceedings of the 8th USENIX conference on Operating systems design and implementation*, ser. OSDI'08. Berkeley, CA, USA: USENIX Association, 2008, pp. 73–86. Available: <http://dl.acm.org/citation.cfm?id=1855741.1855747>
- [67] T. Y. Yeh and G. Reinman, "Fast and fair: data-stream quality of service," in *Proc. CASES '05*. New York, NY, USA: ACM, 2005, pp. 237–248.
- [68] M. Zhang and K. Asanovic, "Victim replication: Maximizing capacity while hiding wire delay in tiled chip multiprocessors," in *Proc. ISCA '05*. Washington, DC, USA: IEEE Computer Society, 2005, pp. 336–345.
- [69] L. Zhao *et al.*, "Towards hybrid last level caches for chip-multiprocessors," *SIGARCH Comput. Archit. News*, vol. 36, no. 2, pp. 56–63, 2008.

A. Response Time Model Design

Alternative Resource Representations

Asynchrony and latency tolerance may make response time components overlap partly or fully; if the latter, then the maximum of the terms might be more appropriate than their sum. The result will still be convex, though, as will any other norm including the 2-norm, *i.e.* the square root of the sum of the squares. This last variation could be viewed as a "partially

overlapped" compromise between the 1-norm (sum) describing no overlap and the ∞ -norm (maximum) describing full overlap.

Sometimes a response time component might be better modeled by a term involving a combination of resources. For example, response time due to memory accesses might be approximated by a combination of memory bandwidth allocation b_{r1} and cache allocation m_{r2} . Such a model could use the geometric mean of the two allocations in the denominator, *viz.* $w_{r1,r2}/\sqrt{b_{r1} \cdot m_{r2}}$, without compromising convexity.

This scheme also accommodates non-bandwidth resources such as memory, the general idea being to roughly approximate "diminishing returns" in the response time with increasing resources. For clarity's sake, rather than using a_r indiscriminately for all allocations, we will denote an allocation of a bandwidth resource by b_r and of a memory resource by m_r . This begs the question of how memory affects the response time. The effect is largely indirect. Memory permits exploitation of temporal locality and thereby *amplifies* associated bandwidths. For example, additional main memory may reduce the need for storage or network bandwidth, and of course increased cache capacity may reduce the need for memory bandwidth. The effectiveness of cache in reducing bandwidth was studied by H. T. Kung [41], who developed tight asymptotic bounds on the bandwidth amplification factor $\alpha(m)$ resulting from a quantity of memory m acting as cache for a variety of computations. He shows that

$$\begin{aligned}\alpha(m) &= \Theta(\sqrt{m}) && \text{for dense linear algebra solvers} \\ &= \Theta(m^{1/d}) && \text{for d-dimensional PDE solvers} \\ &= \Theta(\log m) && \text{for comparison sorting and FFTs} \\ &= \Theta(1) && \text{when temporal locality is absent}\end{aligned}$$

For these expressions to make sense, the argument of α should be dimensionless and greater than 1. Ensuring this might be as simple as letting it be the number of memory resource quanta (*e.g.* hundreds of memory pages) assigned to the application. If a application shows diminishing bandwidth amplification as its memory allocation increases, this can be accommodated:

$$\alpha(m) = \min(c_1 \alpha_1(m), c_2 \alpha_2(m)), \quad c_1, c_2 \geq 0$$

Each bandwidth amplification factor might be described by one of the functions above and included in the denominator of the appropriate component in the response time function model. For example, the storage response time component for the model of an out-of-core sort application might be the quantity of storage accesses divided by the product of the storage bandwidth allocation and $\log m$, the amplification function associated with sorting given a memory allocation of m . Amplification functions for each application might be learned from response time measurements by observing the effect of varying the associated memory resource while keeping the bandwidth allocation constant. Alternatively, redundant components, similar except for amplification function, could be

included in the model to let the model fitting process decide among them.

Initial Model Evaluation

B. Problem Convexity

Penalty Function Convexity

A few facts about convex functions will be useful in what follows. First, a *concave* function is one whose negative is convex. Maximization of a concave function is equivalent to minimization of its convex negative. An affine function, one whose graph is a straight line in two dimensions or a hyperplane in n dimensions, is both convex and concave. A non-negative weighted sum or point-wise maximum (minimum) of convex (concave) functions is convex (concave), as is either kind of function composed with an affine function. The composition of a convex non-decreasing (concave non-increasing) scalar function with a convex function remains convex (concave).

Each penalty function π is the pointwise maximum of two affine functions and is therefore convex. Moreover, since each penalty function is scalar and nondecreasing, its composition with a convex response time function will also be convex.

Response Time Convexity

We now show that response time functions τ including the various bandwidth amplification functions are convex in both the bandwidth and memory resources b_r and m_r given any of the possibilities we have considered. Since norms preserve convexity, this reduces the question to proving each term in the norm is convex. Since all quantities are positive and both maximum and scaling by a positive constant preserve convexity,

$$\begin{aligned} & w / (b \cdot \min(c_1 \alpha_1(m), c_2 \alpha_2(m))) \\ &= \max(w / (b \cdot c_1 \alpha_1(m)), w / (b \cdot c_2 \alpha_2(m))). \end{aligned}$$

It therefore only remains to show that both $1/\sqrt{b \cdot m}$ and $1/(b \cdot \alpha(m))$ are convex in b and m .

A function is defined to be *log-convex* if its logarithm is convex. A log-convex function is itself convex because exponentiation preserves convexity, and the product of log-convex functions is convex because the log of the product is the sum of the logs, each of which is convex by hypothesis. Now $1/b$ is log-convex for $b > 0$ because $-\log b$ is convex on that domain. In a similar way, $\log(1/\sqrt{b \cdot m}) = -(\log b + \log m)/2$ and $\log m^{-1/d} = -(\log m)/d$ are convex. Finally, $\log(1/\log m)$ is convex because its second derivative is positive for $m > 1$:

$$\begin{aligned} \frac{d^2}{dm^2} \log(1/\log m) &= \frac{d^2}{dm^2} (-\log \log m) \\ &= \frac{d}{dm} \left(\frac{-1}{m \log m} \right) \\ &= \frac{1 + \log m}{(m \log m)^2}. \end{aligned}$$

Summing up, a response time function for an application might be modeled by the convex function

$$\begin{aligned} \tau(w, b, \alpha, m) &= \sqrt[p]{\sum_j \left(\frac{w_j}{b_j \cdot \alpha_j(m_j)} \right)^p} \\ &= \|\text{diag} w d^T\|_p \end{aligned}$$

where the w_j are the parameters of the model (the quantities of work) to be learned, the components of d satisfy $d_j = 1/(b_j \cdot \alpha_j(m_j))$, the b_j are the allocations of the bandwidth resources, the α_j are the bandwidth amplification functions (also to be learned), the m_j are the allocations of the memory or cache resources that are responsible for the amplifications. This formulation allows the application response time τ to be modeled as the p -norm of the component-wise product of a vector d that is computed from the resource allocation and a learned vector of work quantities w .

Handling Quasiconvex Response Time Functions

There are examples of response time versus resource behavior that violate convexity. One such example sometimes occurs in memory allocation, where “plateaus” can sometimes be seen as in Figure 10.

Such plateaus are typically caused by algorithm adaptations within the application to accommodate variable resource availability. The response time is really the *minimum* of several convex functions depending on allocation and the point-wise minimum that the application implements fails to preserve convexity. The effect of the plateaus will be a non-convex penalty as shown in Figure 11 and multiple extrema in the optimization problem will be a likely result.

There are several ways to avoid this problem. One is based on the observation that such response time functions will at least be *quasiconvex*. A function f is quasiconvex if all of its *sublevel sets* $S_\ell = \{x | f(x) \leq \ell\}$ are convex sets. Alternatively, f is quasiconvex if its domain is convex and

$$f(\theta x + (1 - \theta)y) \leq \max(f(x), f(y)), 0 \leq \theta \leq 1$$

Quasiconvex optimization can be performed by selecting a threshold ℓ and replacing the objective function with a convex constraint function whose sublevel set S_ℓ is the same as that of f . Next, one determines whether there is a feasible solution for that particular threshold ℓ . Repeated application with a binary search on ℓ will reduce the level of feasibility until the solution is approximated well enough.

Another idea is to use additional constraints to explore convex sub-domains of τ . For example, the affine constraint $a_{p,r} - \mu \leq 0$ excludes application p from any assignment of resource r exceeding μ . Similarly, $\mu - a_{p,r} \leq 0$ excludes the opposite possibility. A binary (or even linear) search of such sub-domains could be used to find the optimal value.

C. Model Update and Downdate Algorithms

C.1. Row Update and Downdate

Downdating makes an instructive example. A row downdate operation applies a sequence of Givens rotations to the rows of Q^T . The rotations are calculated to set every $Q_{i,dd}^T$, $i \neq dd$ to zero. In the end only the diagonal element $Q_{dd,dd}^T$ of column dd will be nonzero. Since Q^T remains orthogonal, the non-diagonal elements of row dd will also have been zeroed automatically and the diagonal element will have absolute value 1. These same rotations are concurrently applied to the elements of $Q^T t$ and to the rows of $R (= Q^T D)$ to reflect the effect that these transformations have on Q^T .

It is crucial to select pairs of rows and an order of rotations that preserves the upper triangular structure of R while zeroing all but the diagonal entry of the chosen column dd of Q^T . Since dd is always below the diagonal of R it initially will contain only zeros. It is therefore sufficient to rotate every non- dd row with row dd , proceeding from bottom to top. The first $m - n - 1$ rotations will keep row $R_{dd,*}$ entirely zero, and the remaining n rotations will introduce nonzeros in $R_{dd,*}$ from right to left. The effect on R will be to replace zero elements by nonzero elements only within row dd . At this point, except for a possible difference in overall sign, $R_{dd,*} = D_{dd,*}$.

Now the rows from 0 down through dd of the modified matrices $Q^T t$ and R and both the rows and columns of the modified Q^T are circularly shifted by one position, moving row dd to the top (and column dd of Q^T to the left edge). The following is the result:

$$\begin{bmatrix} \pm 1 & 0 \\ 0 & \tilde{Q}^T \end{bmatrix} \begin{bmatrix} t_{dd} \\ \tilde{t} \end{bmatrix} = \begin{bmatrix} \pm D_{dd,*} \\ \tilde{R} \end{bmatrix} w - \begin{bmatrix} \pm 1 & 0 \\ 0 & \tilde{Q}^T \end{bmatrix} \begin{bmatrix} \epsilon_{dd} \\ \tilde{\epsilon} \end{bmatrix}$$

The top row has thus been decoupled from the rest of the factorization and may either be deleted or updated with new data.

The update process more or less reverses these steps, adding a new top row to R and t and a row and column to Q^T . Then R is made upper triangular once more by a sequence of Givens rotations that zero its sub-diagonal elements (formerly the diagonal elements of \tilde{R}) one at a time. These rotations are applied not just to R but also to $Q^T t$ and of course to Q^T itself.

C.2. Rank Preservation

Deciding in advance whether downdating a row of R will reduce its rank is equivalent to predicting whether one of the Givens rotations, when applied to R , will zero or nearly zero a diagonal entry of R . This is particularly easy to discover because dd , the row to be downdated, is initially all zeros in R , *i.e.* in the lower part of the matrix. In this situation a diagonal entry of R , $R_{i,i}$ say, will be compromised if and only

if the cosine of the Givens rotation that involves rows dd and i is nearly zero. The result will be an interchange of the zero in $R_{dd,i}$ with the nonzero diagonal element $R_{i,i}$. $R_{dd,i}$ is zero before the rotation because R was originally upper triangular and prior rotations only involved row subscripts greater than i .

ORADIA keeps track of the sequence of values in $Q_{dd,dd}^T$ without actually changing Q^T so that if the downdate at location dd is eventually aborted there is nothing to undo. It is also possible to remember the sines and cosines of the sequence of rotations so they don't have to be recomputed if success ensues. A rank-preserving row to downdate will always be available as long as R is sufficiently "tall". Having at least twice as many rows as columns is enough since the number of available rows to downdate matches or exceeds the maximum possible rank of R .

C.3. Column Update and Downdate

The active-set NNLS method is based on the idea that since the only constraints are variable positivity then for all components either the variable or its gradient will be zero at a solution point; see [9], page 142. The active set, denoted by \mathbf{Z} , comprises the column subscripts j for which the variable w_j is zero and the gradient v_j is positive. If a column j not currently in \mathbf{Z} happens to acquire a negative w_j after a back-solve, w_j is zeroed, j is moved into \mathbf{Z} and column j is downdated in R , thereby making the gradient positive. Conversely, if a column already in \mathbf{Z} happens to acquire a negative gradient v_j it is removed from \mathbf{Z} and updated in R , allowing it to further reduce the value of the objective function.

After initial acquisition of data and QR factorization, each step of ORADIA's NNLS algorithm combines incremental

row and column downdates and updates as follows:

Algorithm C.1: INCREMENTALNNLS(t_0, d_0)

```

local  $R, Q^T, Q^T t, w, v, idx, d, u, done$ 
 $R, Q^T, Q^T t \leftarrow \text{DNDTROW}(R, Q^T, Q^T t, idx)$ 
 $R, Q^T, Q^T t \leftarrow \text{UPDTRW}(t_0, d_0, R, Q^T, Q^T t, idx)$ 
 $w \leftarrow \text{BACKSOLVE}(R, Q^T t, idx)$ 
 $v \leftarrow \text{GRADIENT}(R, Q^T t, idx)$ 
repeat
   $done \leftarrow \text{true}$ 
   $d \leftarrow \arg \min(w)$ 
  if  $w_d < 0$ 
    then
       $done \leftarrow \text{false}$ 
       $R, Q^T, Q^T t, idx \leftarrow \text{DNDTCOL}(R, Q^T, Q^T t, idx, d)$ 
       $w \leftarrow \text{BACKSOLVE}(R, Q^T t, idx)$ 
       $v \leftarrow \text{GRADIENT}(R, Q^T t, idx)$ 
   $u \leftarrow \arg \min(v)$ 
  if  $v_u < 0$ 
    then
       $done \leftarrow \text{false}$ 
       $R, Q^T, Q^T t, idx \leftarrow \text{UPDTCOL}(R, Q^T, Q^T t, idx, u)$ 
       $w \leftarrow \text{BACKSOLVE}(R, Q^T t, idx)$ 
       $v \leftarrow \text{GRADIENT}(R, Q^T t, idx)$ 
until  $done$ 
return  $(w, v)$ 

```

The set \mathbf{Z} and its complement \mathbf{P} are implemented as an index idx containing a vector of the column subscripts comprising \mathbf{P} in increasing order followed by the column subscripts of \mathbf{Z} in increasing order; idx also contains an offset defining the beginning of \mathbf{Z} in the vector. For example, if columns 1, 3, and 4 are in \mathbf{Z} and columns 0, 2, and 5 are in \mathbf{P} then the resulting vector is [0 2 5 1 3 4] and the offset is 3. Since the offset is just the size of the set \mathbf{P} it is naturally called p .

Regardless of status, columns are left in place in R . The columns of R belonging to \mathbf{P} are denoted by R^p and those in \mathbf{Z} by R^z . The updating or downdating of a column only involves modifying the index idx to redefine \mathbf{P} and \mathbf{Z} and then applying Givens rotations to the rows of R to restore R^p to upper triangular form.

When a column indexed by d in R^p is downdated because $w_d < 0$, that column is moved from \mathbf{P} to \mathbf{Z} in idx . To restore R^p to upper triangular form, Givens rotations are applied to R at rows $R_{d,*}$ and $R_{k,*}$ where $d < k < p$. The row subscripts k are used in decreasing order from $p - 1$ down to $d + 1$, and each rotation zeros the subdiagonal element in R^p of the column indexed by k . As usual, these rotations are also applied to Q^T and $Q^T t$. The result in R^z is a “spike” of nonzeros in the column that was moved; it can eventually extend to the bottom of R as row updates occur.

Column movements from \mathbf{Z} to \mathbf{P} are based on the gradient

v of the objective function, namely

$$\begin{aligned}
 v &= 1/2 \nabla \|Dw - t\|_2^2 \\
 &= D^T (Dw - t) \\
 &= R^T Q^T (QRw - t) \\
 &= R^T (Rw - Q^T t) \\
 &= R^T (-Q^T \varepsilon).
 \end{aligned}$$

If for some column in \mathbf{Z} the inner product of the corresponding spiked row in R^T and $-Q^T \varepsilon$ is negative, the column subscript must be moved to \mathbf{P} . Updating R^p reverses the downdating steps by zeroing the spike via a sequence of Givens rotations on R between adjacent pairs of rows, starting at the bottom and ending at $m, m + 1$ where m is the position of the new column in idx . These rotations conveniently extend the columns to the right of m in R^p by one, thus restoring R^p to upper triangular form. Once again, the rotations are also applied to Q^T and $Q^T t$.

A new gradient computation and new back-solve for w are clearly necessary after either downdates or updates to columns of R .